

GCP steps:

GCP steps:

[Steps 0: Create a Github repo](#)

[Step 1: Create a Cloud account](#)

[Step 2: Upload file to Cloud Storage and create service accounts](#)

[Cloud Storage bucket](#)

[Cloud DLP for Cloud Storage](#)

[Optional: Python code to use DLP \(Data Loss prevention as this is sensitive data like gender, email, ip address\)](#)

[Step 3: Python ETL to load data from Cloud Storage to Cloud SQL](#)

[Step 4: Cloud SQL \(MySQL\) database connection for the DEM file](#)

[To import data to a Cloud SQL instance using a CSV file](#)

[Python code to load data from Cloud SQL to BigQuery \(as data will grow, they will need to be in Data Warehouse\)](#)

[Step 5: Terraform code to automate process from Storage to Cloud SQL uploading](#)

[Step 7: ML model - Built on Vertex AI \(Google Cloud\)](#)

[K-Means clustering](#)

[GMM](#)

[Step 8: Deploy the ML model as API endpoint - I used Vertex AI Platform prediction Service](#)

Steps 0: Create a Github repo

<https://github.com/cs-cloud-store/assessment>

Step 1: Create a Cloud account

Google Cloud account with \$300 free trial band set up billing.

Step 2: Upload file to Cloud Storage and create service accounts

Service accounts code:

```
gcloud projects add-iam-policy-binding $GOOGLE_CLOUD_PROJECT \
--member
serviceAccount:https-console-cloud-google-com@hale-brook-377621.iam.gserviceaccount.com
\
--role roles/dlp.admin

gcloud projects add-iam-policy-binding $GOOGLE_CLOUD_PROJECT \
--member
serviceAccount:https-console-cloud-google-com@hale-brook-377621.iam.gserviceaccount.com
\
--role roles/dlp.serviceAgent

gcloud projects add-iam-policy-binding $GOOGLE_CLOUD_PROJECT \
--member serviceAccount:service-`gcloud projects list
--filter="PROJECT_ID:$GOOGLE_CLOUD_PROJECT"
--format="value(PROJECT_NUMBER)">@dlp-api.iam.gserviceaccount.com \
--role roles/viewer
```

Cloud Storage bucket

GSUTIL to have the data uploaded into Storage for project data.

```
gsutil mb gs://dembucket
```

We'll need to create 3 Cloud Storage buckets to store our data:

- Quarantine bucket: our data will initially be uploaded here.
- Sensitive data bucket: the data determined by the DLP API to be sensitive will be moved here.
- Non-sensitive data bucket: the data determined by the DLP API not to be sensitive will be moved here

We can use the gsutil command to create all three of our buckets in one swoop:

#Actual commands to run:

```
gsutil mb gs://dem_quarantine_bucket \
gs://sensitive_dem_bucket \
gs://non_sensitive_dem_bucket
```

```
chandnis3010@cloudshell:~ (hale-brook-377621)$ gsutil mb gs://dem_quarantine_bucket \
gs://sensitive_dem_bucket \
gs://non_sensitive_dem_bucket
Creating gs://dem_quarantine_bucket/...
Creating gs://sensitive_dem_bucket/...
Creating gs://non_sensitive_dem_bucket/...
chandnis3010@cloudshell:~ (hale-brook-377621)$
```

We will use Cloud DLP on the Storage bucket. For future, we can keep 2 copies.

- one sensitive data and
- one non-sensitive data.

Cloud Pub/Sub provides many to many asynchronous messaging between applications. A publisher will create a message and publish it to a feed of messages called a topic. A subscriber will receive these messages by way of a subscription. Based on that subscription, in our case, we'll have a Cloud Function move files to their respective buckets after a DLP job runs.

First, let's create a topic. A message will be published here each time a file is added to our quarantine storage bucket. We'll name it 'classify-topic'

```
gcloud pubsub topics create classify-topic
```

A subscription will be notified when the topic publishes a message. Let's create a pubsub subscription named 'classify-sub':

```
gcloud pubsub subscriptions create classify-sub --topic classify-topic
```

Cloud DLP for Cloud Storage

For now we will be scanning your Cloud Storage buckets with Cloud DLP, and you can set up this new scan job to be run regularly.

- Browse to Cloud Storage in the GCP console, then click on the three-dot menu icon to the right of a relevant bucket. Click on the “Scan with Data Loss Prevention” menu item
- Complete the Cloud DLP scan creation by clicking the “Create” button or, optionally, specify custom configurations such as what info types to inspect for, what sampling options to use, what actions to take, and more.
- Once Cloud DLP scans are completed, you'll get emails with links to the “Scan details” page, where you can analyze findings and take further actions. From there, click on **“View Findings in BigQuery”** to analyze the results.

My First Project ▾

Search (/) for resources, docs, products, and more Search

2 ? C

Job details COPY CANCEL DELETE

demdlpj
Container: gs://dembucket
projects/hale-brook-377621/locations/global/dlpJobs/i-demdlpjjob

Done

OVERVIEW CONFIGURATION

VIEW FINDINGS IN BIGQUERY

Findings	Bytes scanned	Errors
5,960	86.69 KiB	0

Job results

Filter Enter property name or value

Name	Description	Categories	Total	% Total
PERSON_NAME	A full person name, which can include first names, middle names or initials, and last names. Note: Not recommended for use during latency sensitive operations.	Location: GLOBAL Data type: PII	1,694	28.42%
EMAIL_ADDRESS	An email address identifies the mailbox that emails are sent to or from. The maximum length of the domain name is 255 characters, and ...	Location: GLOBAL Data type: PII	1,000	16.78%
IP_ADDRESS	An Internet Protocol (IP) address (either IPv4 or IPv6).	Location: GLOBAL Data type: PII	1,000	16.78%
GENDER	A person's gender identity.	Location: GLOBAL Data type: DEMOGRAPHIC	882	14.8%
LAST_NAME	A last name is defined as the last part of a PERSON_NAME. Note: Not recommended for use during latency sensitive operations.	Location: GLOBAL Data type: PII	668	11.21%
FIRST_NAME	A first name is defined as the first part of a PERSON_NAME. Note: Not recommended for use during latency sensitive operations.	Location: GLOBAL Data type: PII	566	9.5%
LOCATION	A physical address or location. Note: Not recommended for use during latency sensitive operations.	Location: GLOBAL Data type: CONTEXTUAL_INFORMATION	150	2.52%
AGE	An age measured in months or years.	Location: GLOBAL Data type: PII	0	0%
DATE_OF_BIRTH	A date that is identified by context as a date of birth. Note: Not recommended for use during latency sensitive operations.	Location: GLOBAL Data type: PII	0	0%

Explorer + ADD DATA I

Type to search ?

Viewing all resources. [Show starred resources only.](#)

- hale-brook-377621
 - External connections
 - dembucket
 - dipdataset
 - dipitable**

RUN SAVE SHARE SCHEDULE MORE

Q *Unsaved query 2 × ?

1 SELECT * FROM `hale-brook-377621.dipdataset.dipitable` LIMIT 1000

Press Alt+F1 for Accessibility Options

SAVE RESULTS EXPLORE DATA

Query results

JOB INFORMATION RESULTS JSON EXECUTION DETAILS EXECUTION GRAPH PREVIEW

Row	quote	info.type.name	info.type.version	info.type.sensitivity_score.score	likelihood	location...start	location...end	location...start	location...end	location.conte...	co
1	null	GENDER	null	SENSITIVITY_MODERATE	LIKELY	null	null	null	null	gs://dembucket/DE	
2	null	GENDER	null	SENSITIVITY_MODERATE	LIKELY	null	null	null	null	gs://dembucket/DE	
3	null	GENDER	null	SENSITIVITY_MODERATE	LIKELY	null	null	null	null	gs://dembucket/DE	
4	null	GENDER	null	SENSITIVITY_MODERATE	LIKELY	null	null	null	null	gs://dembucket/DE	
5	null	GENDER	null	SENSITIVITY_MODERATE	LIKELY	null	null	null	null	gs://dembucket/DE	
6	null	GENDER	null	SENSITIVITY_MODERATE	LIKELY	null	null	null	null	gs://dembucket/DE	

Load more

Optional: Python code to use DLP (Data Loss prevention as this is sensitive data like gender, email, ip address)

```
import google.cloud.dlp
from google.cloud import dlp_v2
```

```

# Instantiates a DLP client
dlp = dlp_v2.DlpServiceClient()

# Define the info types to be de-identified
info_types = [{"name": "PERSON_NAME"}, {"name": "EMAIL_ADDRESS"}, {"name": "IP_ADDRESS"}]

# Define the de-identification transformation
deidentify_config = {
    "record_transformations": {
        "field_transformations": [
            {
                "fields": [{"name": "first_name"}, {"name": "last_name"}, {"name": "email"}],
                "primitive_transformation": {
                    "character_mask_config": {
                        "masking_character": "*",
                        "number_to_mask": 0.6,
                        "characters_to_ignore": [{"characters_to_skip": ".,:-\'()"}]
                    }
                },
                "fields": [{"name": "ip_address"}],
                "primitive_transformation": {
                    "replace_with_info_type_config": {"info_type": {"name": "IP_ADDRESS"}}
                }
            }
        ],
        "info_type_transformations": {"transformations": [{"info_types": info_types, "action": "redact": {}}]}
    }
}

# Define the re-identification transformation
reidentify_config = {
    "record_transformations": {
        "field_transformations": [
            {
                "fields": [{"name": "first_name"}, {"name": "last_name"}, {"name": "email"}]
            }
        ]
    }
}

```

```

    "primitive_transformation": {"replace_with_info_type_config": {"info_type": {"name": "PERSON_NAME"}},  

        },  

        {"fields": [{"name": "ip_address"}], "primitive_transformation":  

        {"replace_with_info_type_config": {"info_type": {"name": "IP_ADDRESS"}}}},  

        ]  

    },  

    "info_type_transformations": {"transformations": [{"info_types": info_types, "action": {"replaceWithInfoTypeConfig": {"infoType": {"name": "PII"}}}}]},  

}
}

# Define the dataset
dataset = [
    {"id": 1, "first_name": "John", "last_name": "Doe", "email": "john.doe@example.com",
     "gender": "male", "ip_address": "192.168.0.1"},  

    {"id": 2, "first_name": "Jane", "last_name": "Doe", "email": "jane.doe@example.com",
     "gender": "female", "ip_address": "192.168.0.2"},  

]

# De-identify the dataset
deidentify_request = dlp_v2.DeidentifyContentRequest(  

    parent=f"projects/{PROJECT_ID}/locations/{LOCATION}",  

    deidentify_config=deidentify_config,  

    item={"table": {"headers": [{"name": "id"}, {"name": "first_name"}, {"name": "last_name"},  

        {"name": "email"}, {"name": "gender"}, {"name": "ip_address"}], "rows": [{"values":  

            [str(d[key]) for key in ["id", "first_name", "last_name", "email", "gender", "ip_address"]]} for  

            d in dataset]}},  

)
deidentified_response = dlp.deidentify_content(deidentify_request)
deidentified_dataset = [{header["name"]: row["values"][i] for i, header in  

    enumerate(deidentified_response.item.table.headers)} for row in  

    deidentified_response.item.table.rows]

# Re-identify the dataset
reidentify_request = dlp_v2.ReidentifyContentRequest(  

    parent=f"projects/{PROJECT_ID}/locations/{LOCATION}",  

    reidentify_config=reidentify_config,  

    inspect_config={"info_types": info_types},
```

```

item={"table": {"headers": [{"name": "id"}, {"name": "first_name"}, {"name": "last_name"}, {"name": "email"}, {"name": "gender"}, {"name": "ip_address"}], "rows": [{"values": [str(d[key]) for key in ["id", "first_name", "last_name", "email", "gender", "ip_address"]]} for d in deidentified_dataset]}}
)
reidentified_response = dlp.reidentify_content(reidentify_request)
reidentified_dataset = [{header["name"]: row["values"][i] for i, header in enumerate(reidentified_response.item.table.headers)} for row in reidentified_response.item.table.rows]

#print the original and re-identified datasets to test the code
print("Original dataset:")
print(dataset)
print("Re-identified dataset:")
print(reidentified_dataset)
css

```

Step 3: Python ETL to load data from Cloud Storage to Cloud SQL

```

import pandas as pd
import mysql.connector
from google.cloud import storage

# Set up connection to Google Cloud Storage
client = storage.Client()
bucket = client.get_bucket('your-bucket-name')

# Load data from CSV file in GCS into a pandas dataframe
blob = bucket.blob('your-file.csv')
data = blob.download_as_string()
df = pd.read_csv(io.BytesIO(data))

# Transform data (if needed)
# For example, you could concatenate first_name and last_name into a new column called full_name:
df['full_name'] = df['first_name'] + ' ' + df['last_name']

```

```
# Set up connection to MySQL Cloud SQL
cnx = mysql.connector.connect(user='your-db-user', password='your-db-password',
                               host='your-db-host', database='your-db-name')

# Create cursor and table if not exists
cursor = cnx.cursor()
create_table_query = """
CREATE TABLE IF NOT EXISTS customers (
    id INT PRIMARY KEY,
    first_name VARCHAR(50),
    last_name VARCHAR(50),
    email VARCHAR(255),
    gender VARCHAR(10),
    ip_address VARCHAR(50),
    full_name VARCHAR(100)
)
"""

cursor.execute(create_table_query)
cnx.commit()

# Load data into MySQL Cloud SQL
for index, row in df.iterrows():
    insert_query = """
    INSERT INTO customers (id, first_name, last_name, email, gender, ip_address, full_name)
    VALUES (%s, %s, %s, %s, %s, %s, %s)
    """
    values = (row['id'], row['first_name'], row['last_name'], row['email'], row['gender'],
              row['ip_address'], row['full_name'])
    cursor.execute(insert_query, values)
    cnx.commit()

# Close cursor and connection
cursor.close()
cnx.close()
```

Step 4: Cloud SQL (MySQL) database connection for the DEM file

To import data to a Cloud SQL instance using a CSV file

1. Download the file shared by Deloitte as “CSV”
2. Enable “Compute Engine API”
3. Create a development MySQL 8.0 Instance

[← Create a MySQL instance](#)

Instance info

Instance ID * demsql
Use lowercase letters, numbers, and hyphens. Start with a letter.

Password *  GENERATE
Set a password for the root user. [Learn more](#)

No password

▼ PASSWORD POLICY

Database version * MySQL 8.0

▼ SHOW MINOR VERSIONS

Choose a configuration to start with

These suggested configurations will pre-fill this form as a starting point for creating an instance. You can customize as needed later.

- Production
Optimized for the most critical workloads. Highly available, performant, and durable.
- Development
Performant but not highly available, while reducing cost by provisioning less compute and storage.

Production

Development

Id: demsql

Password: demsql

Summary

Region	us-central1 (Iowa)
DB Version	MySQL 8.0
vCPUs	2 vCPU
Memory	8 GB
Storage	100 GB
Network throughput (MB/s) 	500 of 2,000
Disk throughput (MB/s) 	Read: 48.0 of 240.0 Write: 48.0 of 144.0
IOPS 	Read: 3,000 of 15,000 Write: 3,000 of 9,000
Connections	Public IP
Backup	Automated
Availability	Single zone
Point-in-time recovery	Enabled

The screenshot shows the Google Cloud SQL Instances page. At the top, there's a navigation bar with 'Google Cloud' and 'My First Project'. A search bar contains the text 'sql'. On the left, there's a sidebar with 'SQL' and 'Instances' buttons, and links for '+ CREATE INSTANCE' and 'MIGRATE DATA'. The main area has a table with one row for the instance 'demsq1'. The columns include 'Instance ID', 'Type', 'Public IP address', 'Private IP address', 'Instance connection name', 'High availability', 'Location', 'Storage used', 'Labels', and 'Actions'. The instance details are: Type MySQL 8.0, Public IP address 34.6.181.11, Private IP address 10.128.0.3, Instance connection name hale-brook-377621:us-central1:demsq1, High availability ADD, Location us-central1-c, Storage used 0 B of 100 GB, and Labels none.

Python code to load data from Cloud SQL to BigQuery (as data will grow, they will need to be in Data Warehouse)

```
import pandas as pd
import mysql.connector
from google.cloud import bigquery
from google.cloud import storage

# Set up connection to Google Cloud Storage
client = storage.Client()
bucket = client.get_bucket('your-bucket-name')

# Load data from CSV file in GCS into a pandas dataframe
blob = bucket.blob('your-file.csv')
data = blob.download_as_string()
df = pd.read_csv(io.BytesIO(data))

# Transform data (if needed)
# For example, you could concatenate first_name and last_name into a new column called full_name:
df['full_name'] = df['first_name'] + ' ' + df['last_name']

# Set up connection to GCP Cloud SQL
cnx = mysql.connector.connect(user='your-db-user', password='your-db-password',
                               host='your-db-host', database='your-db-name')

# Load data into a pandas dataframe
df = pd.read_sql('SELECT * FROM customers', con=cnx)

# Set up connection to BigQuery
client = bigquery.Client()
table_ref = client.dataset('your_dataset_name').table('your_table_name')

# Create job configuration and load data into BigQuery
```

```

job_config = bq.LoadJobConfig()
job_config.write_disposition = bq.WriteDisposition.WRITE_TRUNCATE
job_config.source_format = bq.SourceFormat.CSV
job_config.schema = [
    bq.SchemaField('id', 'INTEGER'),
    bq.SchemaField('first_name', 'STRING'),
    bq.SchemaField('last_name', 'STRING'),
    bq.SchemaField('email', 'STRING'),
    bq.SchemaField('gender', 'STRING'),
    bq.SchemaField('ip_address', 'STRING'),
    bq.SchemaField('full_name', 'STRING')
]
job = client.load_table_from_dataframe(df, table_ref, job_config=job_config)
job.result() # Wait for job to complete

# Close connection to GCP Cloud SQL
cnx.close()

```

Step 5: Terraform code to automate process from Storage to Cloud SQL uploading

```

# Configure provider
provider "google" {
  project = "hale-brook-377621"
  region  = "your-region"
  zone    = "your-zone"
}

# Configure storage bucket
resource "google_storage_bucket" "data_bucket" {
  name = "your-bucket-name"
}

# Configure Cloud SQL instance
resource "google_sql_database_instance" "db_instance" {
  name        = "your-instance-name"
  region      = "your-region"
}

```

```
database_version = "MYSQL_5_7"
settings {
  tier = "db-f1-micro"
}
}

# Configure Cloud SQL database and user
resource "google_sql_database" "db" {
  name    = "your-db-name"
  instance = google_sql_database_instance.db_instance.name
}

resource "google_sql_user" "db_user" {
  name    = "your-db-user"
  password = "your-db-password"
  instance = google_sql_database_instance.db_instance.name
  host    = "%"
}

# Configure Cloud Function to run ETL code
resource "google_cloudfunctions_function" "etl_function" {
  name      = "your-function-name"
  description = "ETL function that ingests, transforms, and loads data from GCS to Cloud
SQL"

  source_archive_bucket = google_storage_bucket.data_bucket.name
  source_archive_object = "etl_function.zip"

  entry_point = "main"

  runtime = "python37"

  environment_variables = {
    SQL_USER    = google_sql_user.db_user.name
    SQL_PASSWORD = google_sql_user.db_user.password
    SQL_HOST    = google_sql_database_instance.db_instance.ip_address
    SQL_DATABASE = google_sql_database.db.name
    GCS_BUCKET   = google_storage_bucket.data_bucket.name
  }
}
```

```

timeout  = "180s"
available_memory_mb = 256

trigger_http = true
}

# Configure Cloud Scheduler job to run Cloud Function daily
resource "google_cloud_scheduler_job" "etl_job" {
  name      = "your-job-name"
  description = "Schedule ETL function to run daily at 7AM"

  schedule = "0 7 * * *"

  target_type = "google_cloudfunctions_function"
  target_http_method = "POST"
  target_uri = google_cloudfunctions_function.etl_function.https_trigger_url
}

```

Step 6: ETL to load data from Cloud Storage to BigQuery



We will build a data integration pipeline in Cloud Data Fusion for loading, transforming and masking healthcare data in bulk.

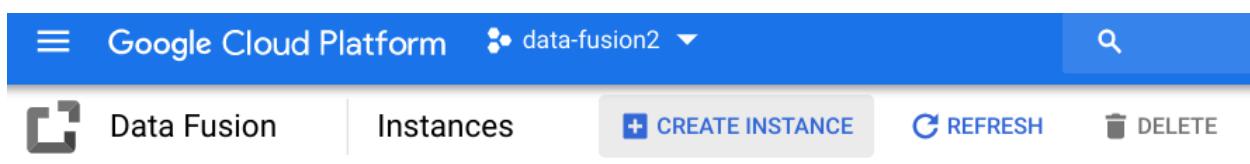
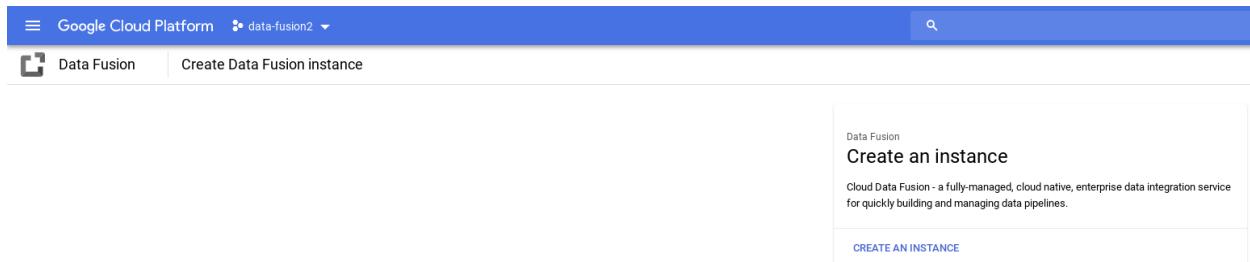
Data Fusion

```

export PROJECT_ID=hale-brook-377621
export BUCKET_NAME=dembucket
export DATASET_ID=demdataset

```

1. Select Data Fusion from the left menu, then click the CREATE AN INSTANCE button in the middle of the page (1st creation), or click the CREATE INSTANCE button at the top menu (additional creation).



3. Provide the instance name. Select Enterprise or Basic.

Google Cloud My First Project

Data Fusion Create Data Fusion instance

Instances Permissions

Instance name * deminstance
Alphanumeric characters, space and - only For eg: My Instance name-1024. Name must start with a letter, 30 character max

Instance ID deminstance

Description

Region us-central1
Region in which the instance is created.

⚠ Selected location is not supported by the zone separation organization policy. [View supported locations](#)

Version 6.7.2

Edition

Developer This edition provides a full-feature edition for product exploration and development environments with zonal availability and limitation on execution environment.

Basic This edition provides comprehensive data integration capabilities. Users can build batch data pipelines; connect to any data source; perform code-free transformations. Limitation on simultaneous pipeline runs. Recommended for non-critical environments.

Enterprise This edition provides all the functionality provided in the Basic edition. In addition, includes support for realtime data pipelines; interactions with data lineage; higher scalability; and high availability. Recommended for critical environments.

ADD ACCELERATORS 0 accelerators added

Authorization

⚠ Cloud Data Fusion does not have permission to use the selected Dataproc Service Account. These permissions are required to launch Dataproc clusters in customer projects for executing pipelines. It is necessary to grant permission before the Cloud Data Fusion instance can be created.

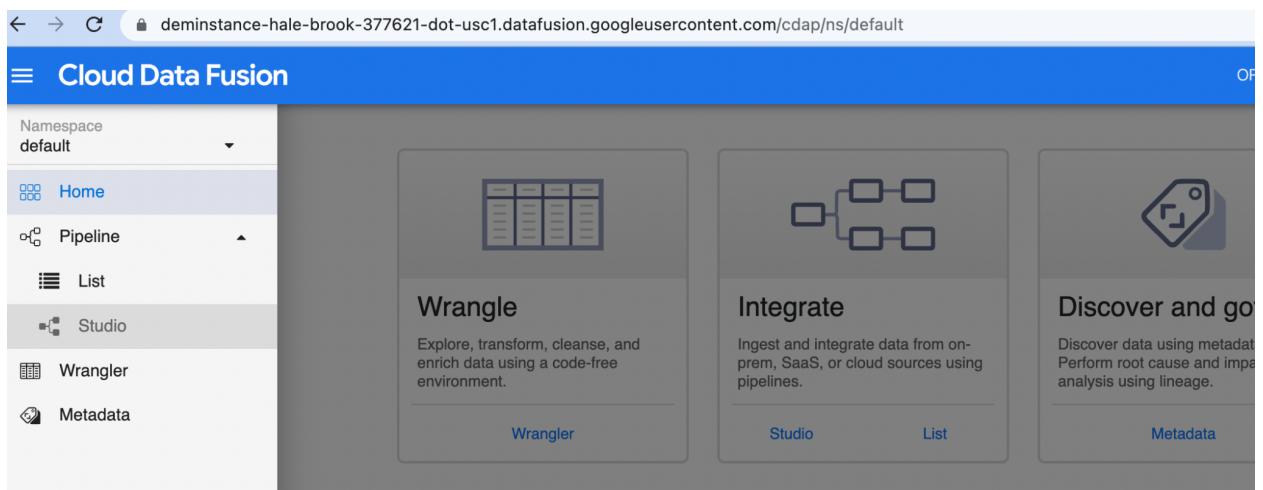
GRANT PERMISSION **MORE INFO**

Dataproc Service Account 604147830502-compute@developer.gserviceaccount.com
A service account is used to authorize the instance to run a Dataproc job.

- 4.
5. Copy the service account.

6. On the IAM permissions page, we will now add the service account as a new member and grant it the Cloud Data Fusion API Service Agent role. Click the Add button, then paste the "service account" in the New members field and select Service Management -> Cloud Data Fusion API Server Agent role.
7. Check for default SSH in VPC. If not there, then create:

```
gcloud beta compute
--project=hale-brook-377621 firewall-rules create default-allow-ssh --direction=INGRESS
--priority=1000 --network=default --action=ALLOW --rules=tcp:22
--source-ranges=0.0.0.0/0 --enable-logging
```
8. Now that we have the Cloud Fusion environment in GCP let's build a schema. We need this schema for transformation of the CSV data. In the Cloud Data Fusion window, click the View Instance link in the Action column. You will be redirected to another page. Click the provided url to open Cloud Data Fusion instance. Your choice to click "Start Tour" or "No, Thanks" button at the Welcome popup. Expand the "hamburger" menu, select Pipeline -> Studio



9. By default, the initial import will assume there is only one column in your data file. To parse it as a CSV, choose Parse → CSV, then select the delimiter and check the "Set first row as header" box as appropriate. Click the Apply button.

Cloud Data Fusion | Studio

WRANGLE

Select Data

Root / dembucket

Format * csv

Enable Quoted Values False

Use First Row as Header True

File encoding UTF-8

IMPORT SCHEMA

CANCEL CONFIRM

deminstance-hale-brook-377621-dot-usc1.datafusion.googleusercontent.com/pipelines/ns/default/studio

Cloud Data Fusion | Studio

WRANGLE

DEM_Challenge_Section1_DATASET.xlsx - DAT

	Int	String	String	String	String	String	String
	# id	first_name	last_name	email	gender	ip_address	
1	1	Margareta	Lughtisse	mlaughtisse0@mediafire.com	Genderfluid	34.148.232.131	
2	2	Vally	Garment	vgarment1@wisc.edu	Bigender	15.158.123.36	
3	3	Tessa	Curee	tcurree2@php.net	Bigender	132.209.143.225	
4	4	Arman	Heineking	aheineking3@tuttocitta.it	Male	157.110.61.233	
5	5	Rosella	Trustie	rtrustie4@ft.com	Non-binary	49.55.218.81	
6	6	Roxie	Springett	rspringett5@deviantart.com	Male	51.206.104.138	
7	7	Gabi	Kernell	gkernell6@hugedomains.com	Female	223.30.27.146	
8	8	Dino	Kentwell	dkentwell7@com.com	Agender	107.244.52.181	
9	9	Petronilla	Jandel	pjandel8@amazon.co.uk	Female	187.54.208.203	
10	10	Courtnay	Zecchinelli	czechinelli9@cam.ac.uk	Genderfluid	80.96.245.191	
11	11	Sunny	Kennermann	skennermann@quantcast.com	Genderqueer	211.13.246.106	
12	12	Dayle	McCrachen	dmccrachenb@booking.com	Genderqueer	159.232.55.236	

Columns (6) Transformation steps (0)

Search Column names

Apply More

- The "Columns" and "Transformation steps" tabs show output schema and the Wrangler's recipe. Click Apply at the upper right corner. Click the Validate button. The green "No errors found" indicates success.

11. In Wrangler Properties, click the Actions dropdown to Export the desired schema into your local storage for future Import if needed.
12. Save the Wrangler Recipe for future usage.
13. In the Data Pipelines UI, in the upper left, you should see that Data Pipeline - Batch is selected as the pipeline type.
14. Select the Source node.
15. Under the Source section in the Plugin palette on the left, double-click on the Google Cloud Storage node, which appears in the Data Pipelines UI.
16. Point to the GCS source node and click Properties.
17. Source Node: Fill in the required fields. Set following fields:
 - Label = {any text}
 - Reference name = {any text}
 - Project ID = auto detect
 - Path = GCS URL to bucket in your current project. For example, gs://\$BUCKET_NAME/csv/
 - Format = text
 - Path Field = filename
 - Path Filename Only = true
 - Read Files Recursively = true
 - Add field 'filename' to the GCS Output Schema by clicking the + button.
 - Click Get Schema to generate schema.
16. Select the Transform node.
 - Under the Transform section in the Plugin palette on the left, double-click the Wrangler node, which appears in the Data Pipelines UI. Connect GCS source node to Wrangler transform node.
 - Point to the Wrangler node and click Properties.
 - Click Actions drop down and select Import to import a saved schema
17. Data masking and de-identification
 - You can select individual data columns by clicking the down arrow in the column and applying masking rules under the Mask data selection as per your requirements (for example, Gender and email address column).

- I masked the gender as a test to show that the function works. I selected Show last 2 characters only.

The screenshot shows the Google Cloud Data Studio Wrangle interface. A context menu is open over the 'gender' column, specifically at row 10 (Courtney). The menu path 'Mask data' is highlighted. Below it, the option 'Show last 4 characters only' is selected, with a value of '91' shown. Other options like 'Show last 2 characters only' (value 106) and 'Custom selection' are also visible. The main table view shows 13 rows of sample data from a 'DEM_Challenge_Section1_DATASET.xlsx' file.

18. Sink node: Select the sink node.

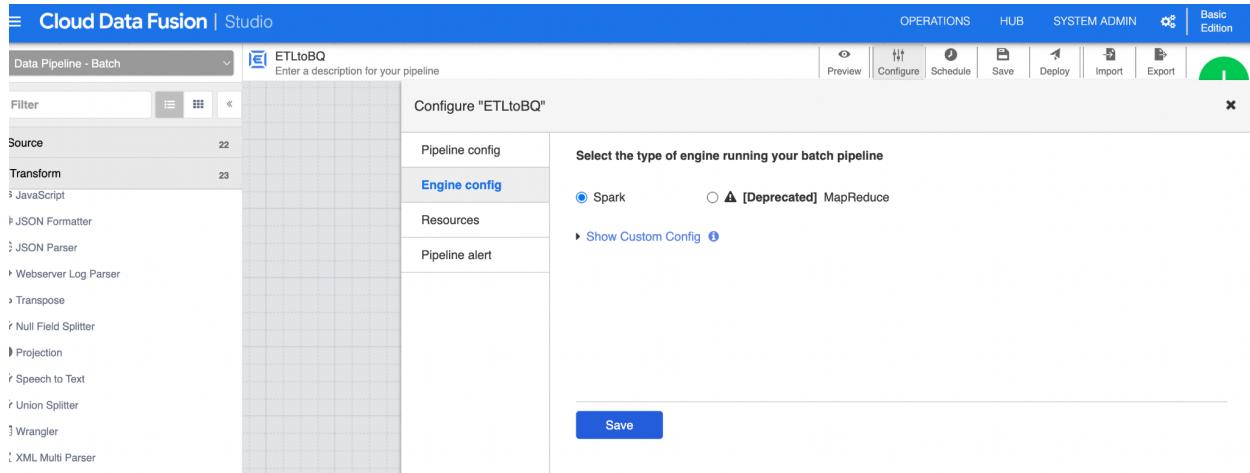
- Under the Sink section in the Plugin palette on the left, double click on BigQuery node, which will appear in the Data Pipeline UI.
- Point to the BigQuery sink node and click Properties.
- Drag a connection arrow > on the right edge of the source node and drop on the left edge of the destination node.

The screenshot shows the Cloud Data Fusion Studio BigQuery Properties page for version 0.20.3. The 'Basic' tab is selected. Key configuration fields include:

- Reference Name:** datafusionref
- Dataset:** demdataset
- Table:** demtable
- Temporary Bucket Name:** Google Cloud Storage bucket for temporary data
- GCS Upload Request Chunk Size:** GCS upload request chunk size in bytes
- Operation:** Radio buttons for Insert (selected), Update, and Upsert.

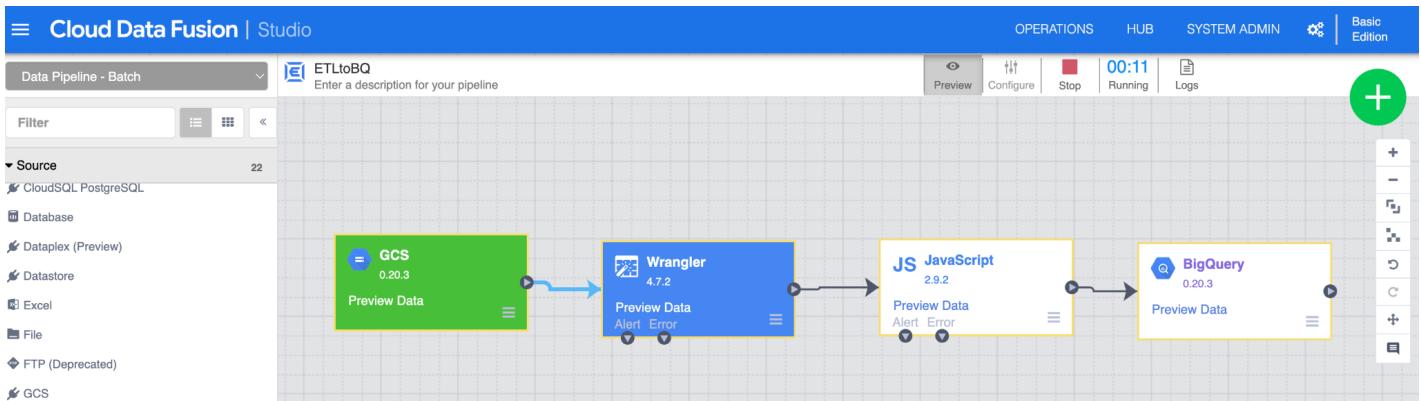
The top bar shows 'Cloud Data Fusion | Studio' and 'Basic Edition'. The status bar indicates 'No errors found.' and has a 'Validate' button.


19. Automate: In the upper-right corner of Data Fusion studio, click Configure. Select Spark for Engine Config. Click Save in Configure window.



Note: When you run a pipeline, Cloud Data Fusion provides an ephemeral Cloud Dataproc cluster, runs the pipeline, and then tears down the cluster. This could take a few minutes. You can observe the status of the pipeline transition from Provisioning to Starting to Running to Deprovisioning to Succeeded during this time.

Running the job without errors and successfully !!!





Step 7: ML model - Built on Vertex AI (Google Cloud)

Perform K-Means clustering and GMM clustering as the data looks like demographics data of people and contains network information (IP address). My thought process is that we can find patterns and outliers from the people using clustering algorithms. These are unsupervised type of machine learning.

K-Means clustering

```
from google.cloud import aiplatform  
from google.protobuf import json_format
```

```
from google.protobuf.struct_pb2 import Value
from google.cloud.aiplatform_v1.types import (
    DeployedModel,
    ExplanationMetadata,
    ExplanationParameters,
    ExplanationSpec,
    InputDataConfig,
    Model,
    ModelContainerSpec,
    Port,
    PredictRequest,
    PredictionInput,
    PredictionOutput,
)
import pandas as pd
from sklearn.preprocessing import StandardScaler
import numpy as np

# Set the project and location
PROJECT_ID = "hale-brook-377621"
LOCATION = "us-central1"

# Set the name and location of the dataset file
DATASET_PATH = "gs://your-bucket/dataset.csv"

# Create a Vertex AI client
client_options = {"api_endpoint": f"{LOCATION}-aiplatform.googleapis.com"}
client = aiplatform.gapic.JobServiceClient(client_options=client_options)

# Load the dataset
df = pd.read_csv(DATASET_PATH)

# Select features for clustering
X = df[["first_name", "last_name", "email", "gender", "ip_address"]]

# Preprocess the features using standard scaler
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
```

```
# Set the number of clusters
num_clusters = 3

# Set the container image URI
image_uri = "us-docker.pkg.dev/vertex-ai/prediction/kmeans-cpu.1-0:latest"

# Create the model container specification
model_container_spec = ModelContainerSpec(image_uri=image_uri)

# Create the K-Means model
model = Model(
    display_name="kmeans-model",
    container_spec=model_container_spec,
    metadata_schema_uri="gs://google-cloud-aiplatform/schema/prediction/classification_
1.0.0.yaml",
)

# Create the input data configuration
input_data_config = InputDataConfig(
    instances_format="json",

    json_instance_schema_uri="gs://google-cloud-aiplatform/schema/prediction/input/text
_classification_example.json",
)

# Create the explanation spec
explanation_metadata = ExplanationMetadata(
    inputs={"features": {"input_tensor_name": "input"}}
)
explanation_parameters = ExplanationParameters(
    {"sampled_shapley_attribution": {"path_count": 20}}
)
explanation_spec = ExplanationSpec(
    metadata=explanation_metadata, parameters=explanation_parameters
)

# Create the model deployment
deployed_model = DeployedModel(model=model)
```

```

endpoint = aiplatform.Endpoint.create(
    display_name="kmeans-endpoint", explanation_spec=explanation_spec
)
endpoint.deploy(
    deployed_model=deployed_model,
    traffic_percentage=100,
    machine_type="n1-standard-2",
)

# Create a prediction request for a sample dataset
input_data = [{"first_name": row[0], "last_name": row[1], "email": row[2], "gender": row[3],
"ip_address": row[4]} for row in X.values]
predict_request = PredictRequest(
    endpoint=endpoint.resource_name,

instances=[json_format.Parse(json_format.MessageToJson(PredictionInput(data=Input
DataConfig.EncodedData(json_format.Parse(json_format.MessageToJson(Value(json_v
alue=bytes(json.dumps(instance), "utf-8")))))), PredictRequest) for instance in
input_data],
)

# Make a prediction request
response = client.predict(predict_request)

# Get the cluster assignments
cluster_assignments = np.array([np.argmax(prediction.outputs["scores"].value) for
prediction in response.predictions])

#add a new column in the dataset
df["cluster"] = cluster_assignments
#print the dataset with the predicted clusters
print(df.head())

```

GMM

```

from google.cloud import aiplatform
from google.cloud.aiplatform.gapic.schema import predict as predict_pb2
import numpy as np

```

```
import pandas as pd
from sklearn.preprocessing import StandardScaler
from sklearn.mixture import GaussianMixture

# Load the dataset
df = pd.read_csv("dataset.csv")

# Select the features for clustering
X = df[["first_name", "last_name", "email", "gender", "ip_address"]]

# Preprocess the features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Set the number of clusters
num_clusters = 3

# Set the container image URI for the GMM model
container_image_uri =
"us-docker.pkg.dev/vertex-ai/prediction/gaussian-mixture-model:latest"

# Create the model container specification
model_container_spec = {
    "image_uri": container_image_uri,
}

# Create the GMM model
model = aiplatform.Model(
    display_name="gmm-clustering-model",
    container_spec=model_container_spec,
    predict_schemata=predict_pb2.ValueSpec(
        encoded_value=predict_pb2.EncodedValueSpec(
            tensor_shape=predict_pb2.TensorShape(dim=[predict_pb2.FixedDim(size=5)])
        )
    ),
)

# Create the input data configuration
input_data_config = {
```

```

"instances": [
    {
        "first_name": x[0],
        "last_name": x[1],
        "email": x[2],
        "gender": x[3],
        "ip_address": x[4],
    }
    for x in X_scaled
]
}

# Set the number of clusters in the explanation metadata
explanation_metadata = {
    "inputs": {
        "first_name": {"type": "CATEGORY"},
        "last_name": {"type": "CATEGORY"},
        "email": {"type": "CATEGORY"},
        "gender": {"type": "CATEGORY"},
        "ip_address": {"type": "CATEGORY"},
    },
    "outputs": {"cluster": {"type": "CATEGORY", "number_of_categories": num_clusters}},
}
}

# Set the explanation spec
explanation_spec = {"metadata": explanation_metadata}

# Train the GMM model and deploy it to an endpoint
endpoint = model.deploy(
    machine_type="n1-standard-4",
    min_replica_count=1,
    max_replica_count=1,
)
)

# Create a prediction request for a sample dataset
sample_dataset = [
    {"first_name": "Alice", "last_name": "Smith", "email": "alice@example.com", "gender": "Female", "ip_address": "123.45.67.89"},


```

```

        {"first_name": "Bob", "last_name": "Jones", "email": "bob@example.com", "gender": "Male", "ip_address": "98.76.54.32"},  

        {"first_name": "Charlie", "last_name": "Lee", "email": "charlie@example.com", "gender": "Male", "ip_address": "12.34.56.78"},  

    ]  

sample_dataset_scaled =  

scaler.transform(pd.DataFrame(sample_dataset)[["first_name", "last_name", "email",  

"gender", "ip_address"]])  

prediction_request = predict_pb2.PredictRequest(  

    instances=[{"values": [list(row)]} for row in sample_dataset_scaled],  

    parameters={"explanation_spec": explanation_spec},  

)  

# Make the prediction request using the Vertex AI Python SDK  

response = endpoint.predict(prediction_request)  

# Get the cluster assignments from  

# Extract the predicted clusters  

predicted_clusters = np.array(response.predictions[0].tables.values).flatten()  

# Print the predicted clusters for each sample dataset  

for i, cluster in enumerate(predicted_clusters):  

    print(f"Sample dataset {i+1} is assigned to cluster {cluster}")

```

Step 8: Deploy the ML model as API endpoint - I used Vertex AI Platform prediction Service

```

from googleapiclient import discovery
from google.oauth2 import service_account
import numpy as np

# Specify the project ID and the name of the model to be deployed
PROJECT_ID = 'hale-brook-377621'
MODEL_NAME = 'your-model-name'

```

```

# Load the service account credentials
credentials =
service_account.Credentials.from_service_account_file('path/to/your/credentials.json')

# Create a client object for the AI Platform Prediction service
ml = discovery.build('ml', 'v1', credentials=credentials)

# Define a function to preprocess the input data before making predictions
def preprocess(input_data):
    # Apply any necessary preprocessing steps here, such as scaling or one-hot encoding
    return input_data

# Define a function to make predictions using the deployed model
def predict(input_data):
    # Preprocess the input data
    preprocessed_data = preprocess(input_data)

    # Create a request body for the prediction request
    request_body = {"instances": preprocessed_data.tolist()}

    # Make the prediction request using the AI Platform Prediction API
    parent = f'projects/{PROJECT_ID}/models/{MODEL_NAME}'
    response = ml.projects().predict(name=parent, body=request_body).execute()

    # Extract the predicted labels from the response
    predicted_labels = np.array(response['predictions']).flatten()

    # Return the predicted labels
    return predicted_labels

```

```

gcloud ai-platform models create your-model-name --regions=us-central1
gcloud ai-platform versions create your-version-name --model=your-model-name
--runtime-version=2.5 --python-version=3.7 --framework=scikit-learn
--origin=gs://your-bucket/path/to/model --project=hale-brook-377621

```